

Gaps in Operator Fusion for Deep Learning Systems: The Trade-off Between Compilation and Execution of Unfused Operations

Leonardo Yukio Macedo Kamei
PUC-MG
Belo Horizonte, Brazil
leokamei1@gmail.com

Pedro Henrique Ramos Costa
PUC-MG
Belo Horizonte, Brazil
phramoscosta@gmail.com

Abstract

Operator fusion can reduce steady-state latency, but compiler work becomes relevant in short-lived inference workloads. We characterize native TVM, XLA, and PyTorch 2 (*TorchInductor*) stacks on ResNet-18/50, BERT, and GPT-2 using $K = 5$ cold processes per model and input, 10 warmups, and 50 timed executions per process. We use the amortized cost model $T(n) = T_C + n T_X$, where T_C is the backend-specific compilation-cost proxy, T_X the steady-state latency, and n the number of executions, to expose the regimes favored by each measured stack. For TorchInductor, an inference-only *Conv-BatchNorm* parameter fold applied as offline preprocessing reduced the mean compiler-cost proxy by 29.5–41.1% for ResNet-18 and 1.9–28.6% for ResNet-50; nine of ten differences remain Holm-resolved after adding the measured 100–120-ms preprocessing cost. Execution effects were smaller and model-dependent. An audit of a previous BERT experiment found that its matcher performed zero rewrites and that the apparent gain came from compiling an unchanged graph second in the same process; we therefore withdraw that result. The benchmark reports measured system-level trade-offs, while neither the native-stack comparison nor the valid ResNet fold establishes a universal compiler pass ordering.

Keywords

Operator Fusion, Machine Learning Compilers, Compilation Time, Kernel Fusion, Performance Benchmarking

1 Introduction

The *operator fusion* technique has direct roots in classical compiler optimizations, such as *loop* fusion, historically used to increase parallelism and improve data locality [12]. In modern compilers for Machine Learning, particularly on GPUs, this technique plays a central role: by combining multiple operations into a single *kernel*, memory traffic is reduced, intermediate tensors are avoided, and the number of dynamic launches during execution is decreased [4, 8]. Systems such as XLA, TVM, TensorRT, and the TorchInductor backend adopt aggressive fusion to amortize CPU–GPU communication costs and consolidate operator sequences into larger, more efficient *kernels* [1, 4, 7].

Fusion is performed over computational/dataflow graphs or compiler intermediate representations. Deciding *what* and *how* to fuse is challenging: specific graph-partitioning formulations of operator fusion can be mapped to NP-complete problems, while classical loop-fusion formulations are likewise computationally difficult [3, 5]. For this reason, compilers use legality rules, heuristics, and cost models. Compilation cost is especially important in dynamic

workloads and scenarios with few repetitions, where fixed compiler work can dominate perceived performance.

Many compiler evaluations emphasize steady-state latency, even though deployment decisions may also depend on compilation cost. Reporting both quantities under repeated cold processes makes explicit when additional fixed cost can be amortized. XLA [11], TVM [4], and TorchInductor (PyTorch 2) [1] use different frontend contracts, optimization pipelines, and code-generation strategies, resulting in different compilation- and execution-time profiles.

We evaluate this trade-off for native TVM, XLA, and TorchInductor ResNet implementations under a common outer sampling protocol. For n executions, we use $T_b(n) = a_b n + b_b = T_C + n T_X$, where $b_b = T_C$ is the compilation-cost proxy and $a_b = T_X$ is steady-state latency (Section 4). Minimizing the measured lines gives a descriptive deployment envelope for this hardware and software stack. Because the backends use native model implementations, layouts, parameter-binding contracts, and compilation APIs, the comparison is system-level rather than an isolation of compiler quality. We additionally evaluate a legal *Conv-BatchNorm* fold as offline TorchInductor preprocessing.

Contributions:

- A repeated cold-process protocol and the amortized model $T(n) = T_C + n T_X$ for reporting compilation–execution trade-offs;
- A $K = 5$ system-level characterization of native XLA, TorchInductor, and TVM stacks on ResNet-18/50, BERT, and GPT-2;
- Artifact analysis (TVMScript, HLO, and Triton) and a correctness-audited TorchInductor *Conv-BatchNorm* fold, including the withdrawal of an inapplicable BERT fold result.

2 Related Work

Operator fusion can be seen as the computation-graph analogue of classical *peephole optimization*. Fraser’s machine-independent peephole optimizer characterized valid local rewrites over a sequence of machine instructions by describing each instruction’s effect as a Register Transfer List, replacing a window of instructions by a single equivalent one whenever such a description existed [6]. Operator fusion plays the same role one level up: where peephole optimization rewrites a window of three-address-code instructions into an equivalent instruction, fusion rewrites a window of computation-graph operators into an equivalent operator that preserves the observable semantics.

Among graph-level approaches, *Optimus* minimizes memory accesses over a computation DAG through an offline dynamic-programming algorithm and an online buffer-budget variant [2]. DNNFusion combines data-flow analysis with fusion templates and a cost model, illustrating that legality and profitability are distinct concerns [8]. The XNNC study maps its directed-graph partitioning formulation to an NP-complete problem and reports that a backtracking solver did not finish the largest evaluated networks within a one-hour limit, whereas its greedy clustering alternatives terminated on the full suite [3]. Production compilers similarly combine restricted search, pattern rules, and backend-specific code generation [1, 4].

3 Theoretical Background

Consider a sequence of operators $x_{k+1} = f_k(x_k)$ and $x_{k+2} = f_{k+1}(x_{k+1})$, in which each f_i acts on tensors of a computation graph. We say that f_k and f_{k+1} are *fusible* when there exists a single operator h such that $h(x_k) = f_{k+1}(f_k(x_k))$ and the execution of h preserves every observable semantics of the original sequence. In this case, the compiler performs the transformation $(f_k, f_{k+1}) \Rightarrow h$, eliminating the materialization of intermediate tensors, reducing the number of kernel launches, and improving memory locality [4, 8].

As discussed in Section 2, this notion of fusibility is the computation-graph analogue of classical *peephole optimization*: h plays, for a window of graph operators, the role that an equivalent instruction plays for a window of three-address-code instructions. In ML models, however, operators handle multi-dimensional tensors and their definitions are often complex. We therefore distinguish kernel fusion from parameter folding. Equations (III) and (VIII) in Appendix A describe BatchNorm+ReLU fusion and inference-only Conv--BatchNorm reparameterization followed by ReLU. Equation (X) gives a conditional affine LayerNorm--Linear reparameterization; our BERT audit found no legal instance of this last pattern.

Kernel sequences may be left unfused because of legality, profitability, or backend restrictions [1, 8]:

- **Training mode.** A forward-only BatchNorm parameter fold is invalid in training because batch statistics and state updates remain dynamic. Training-time kernel fusion remains possible when an implementation preserves the required backward values and state semantics.
- **Graph uses.** A local reparameterization must account for every consumer of the rewritten value. Multiple consumers do not universally forbid kernel fusion, but they can make a single-consumer fold illegal.
- **Numerical precision.** Mixed-precision rewrites must preserve conversion and accumulation semantics within the accepted numerical tolerance.
- **Shapes and indexing.** Broadcasting, layout conversion, reindexing, or dynamic shapes can make a candidate illegal or unprofitable.
- **Operation type.** Differing memory-access patterns or synchronization requirements may prevent profitable execution in one generated kernel.

- **Backend restrictions.** External-library calls such as cuDNN or cuBLAS constrain which surrounding operations a compiler can incorporate into the same call [1, 11].

In every case, the fusion decision involves an explicit trade-off: investing more compilation time to reduce execution time, or compiling quickly with fewer fusions. This balance underlies the cost model adopted in this work and guides the comparison between backends.

4 Methodology

We compare *TVM* [4], *XLA* [11], and *TorchInductor* (PyTorch 2) [1] in a controlled environment, measuring **compilation time**, **average latency per execution**, and **fusion granularity**.

4.1 Experimental Environment

The repeated campaigns were run on Fedora 41, an AMD Ryzen 5 4600G CPU, 15.4 GiB of RAM, and an NVIDIA RTX 3050 with 8 GB of VRAM, using driver 580.105.08. The host *nvcc* compiler was CUDA 12.6 (V12.6.77). Transformer measurements ran in the archived container on host kernel 6.17.10; every JSON records the effective Python, framework, and platform versions.

Per-framework configuration. For the **TVM** experiments, we used Python 3.11.13, cuDNN version 9.1 (via PyTorch), TVM 0.22.dev0 (commit 3b60f1c), LLVM 20.1.8, and CUDA 12.6 integrated into the TVM backend.

For **XLA/TorchInductor**, the ResNet environment used Python 3.10.18. XLA used JAX/jaxlib 0.6.2 and PJRT (jax-cuda1 2-pjrt 0.6.2). Its CUDA plugin was built against cuDNN 9.8 and explicitly loaded the host cuDNN 9.11.0 at runtime; the packaged cuDNN 9.10.2 had produced intermittent convolution-plan failures on this machine. The repeated ResNet campaign uses XLA’s default autotuner, records both build/runtime cuDNN versions, and retains a failed process only as retry metadata, never as a timing observation. The Transformer container uses Python 3.12.3 and Transformers 4.56.2: TorchInductor uses PyTorch 2.9.0+cu128, XLA uses JAX 0.6.2, and TVM 0.22.dev0 imports the PyTorch 2.5.1+cu124 graph. The BERT fold-control audit in Table 6 uses the same pinned TorchInductor environment. Because compiler and model-library versions affect compilation cost and generated graphs, absolute times are compared only within the corresponding campaign and table.

The common outer conditions are FP32, the same GPU, inference mode, 10 warmups, 50 timed executions, explicit synchronization, and isolated cold processes. Backend-native layouts are retained: TorchInductor uses logical NCHW tensors in `channels_last` memory format, XLA uses NHWC, and TVM uses NCHW. The TVM environment is isolated at version 0.22.dev0.

4.2 Models and Inputs

We evaluated **ResNet-18/50**, **BERT**, and **GPT-2**. Each ResNet input is written as (N, C, H, W) : batch size N , number of channels C , and spatial height and width $H \times W$. Inputs for ResNets (CNNs): $(1, 3, 224, 224)$, $(16, 3, 224, 224)$, $(64, 3, 224, 224)$, $(1, 3, 512, 512)$, $(1, 3, 1024, 1024)$. BERT uses $(\text{batch}, \text{seq_len})$ pairs $(1, 64)$, $(1,$

128), and (8, 128); GPT-2 uses (1, 16), (1, 128), (8, 128), and (8, 256).

TorchInductor and TVM use the same nominal torchvision ResNet definitions with `weights=None` and seed 0; the original and folded TorchInductor children start from identical state dictionaries, while TVM imports the PyTorch FX graph and binds its parameters. XLA uses the native flaxmodels ResNet definitions with `pret_rained=None`, no embedded input normalization, logits output, PRNG seed 0, and variables closed over the JIT-compiled function. Inputs are seeded random tensors for TorchInductor/TVM and ones for XLA. Consequently, the experiment compares complete native software stacks for the same named architecture, shape, precision, and hardware; it does not establish numerical model equivalence or isolate the compiler while holding frontend graph, parameters, layout, and parameter-binding semantics constant.

For Transformers, TorchInductor and TVM start from the same Hugging Face base configurations, deterministic random weights, and input IDs; XLA uses the corresponding native Flax model. GPT-2 receives a fixed four-dimensional causal mask, avoiding a nested v map masking path that PyTorch 2.5 cannot export. For TVM import only, the exported `aten._unsafe_view` alias is normalized to the equivalent supported `aten.reshape`; no tensor values or shapes change.

The **ResNet-18/50**, **BERT**, and **GPT-2** models were chosen because they represent distinct and complementary architectural patterns. The ResNets characterize classical convolutional architectures, with recurring *Conv–BatchNorm–ReLU* chains, which are canonical cases for analyzing operator fusion in CNNs.

BERT represents *Transformer* architectures based on encoders, dominated by matrix multiplications, linear projections, normalizations, and activations, allowing us to evaluate fusion strategies on attention-centric graphs.

GPT-2, in turn, exemplifies autoregressive *decoder-only* models, with a more linear and repetitive computational flow, typical of generative models, in which fusion granularity directly affects inference latency.

Together, these models allow us to analyze the behavior of fusion heuristics across different computational and architectural regimes.

4.3 Experimental Pipeline

The repeated workflow adopts six outer stages shared by TVM, XLA, and TorchInductor for both convolutional and Transformer models.

(1) *Compilation*. For each (*model*, *input*, *backend*) configuration, we perform $K = 5$ cold compilations in independent processes. TorchInductor’s original and folded ResNet variants are also compiled in separate processes. TVM and XLA children compile only the optimized/JIT variant reported in the comparison, so an auxiliary variant cannot warm in-process compiler state first. Each process receives private TorchInductor, Triton, and JAX compiler-cache directories; persistent TorchInductor and JAX compilation caches are disabled. Model-weight files and the operating system page cache remain shared, and model loading is outside the timed interval. We report mean, sample standard deviation, and two-sided 95% confidence intervals over the five process means. The independently compiled process—not an execution from an already

compiled process—is the experimental unit. The separate BERT fold audit follows the same $K = 5$ cold-process protocol.

(2) *Warmup and execution + graph capture*. Execution occurs in two phases: (i) a *warmup* phase for runtime state and *auto-tuning*; and (ii) synchronized timed executions. The benchmark additionally exports backend IR artifacts—optimized HLO (XLA), TVMScript/TIR (TVM), and generated Triton/Python wrappers (TorchInductor)—outside the steady-state timing window and collects structural call metadata when available.

(3) *Analysis of kernel quantity and quality*. From the captured IRs, we perform a structural analysis of the *kernels* generated by each backend. This stage includes quantifying the total number of *kernels*, distinguishing between internal *kernels* and external library calls, as well as qualitatively characterizing the fusions performed. This analysis allows us to evaluate fusion granularity and the degree of delegation to optimized libraries, providing structural context for interpreting the performance results.

(4) *Measurements and algebraic model*. The timing metrics are a backend-specific compilation-cost proxy and the average synchronized latency T_X . In ResNet, each of the $K = 5$ isolated processes performs 10 warmups followed by 50 timed executions. Execution confidence intervals use Student’s t distribution over the five process means; the 250 individual samples are retained as raw observations but are not treated as 250 independent compilations. For **TVM**, the proxy times graph passes and `relax.build` after FX tracing/import. For **TorchInductor** and **XLA**, it is estimated by

$$T_C = T_{\text{first}} - T_X,$$

where T_{first} aggregates the backend’s first compiled execution and T_X is steady-state latency. The estimator is applied within each process before aggregation. These operational windows follow the available backend APIs: they share the cold-process sampling design but do not include identical frontend work. Cross-backend T_C values and envelopes must therefore be interpreted as native-stack measurements under this harness, not as an isonomic comparison of compiler passes alone.

Algebraic model. We use the following amortized model for compilation followed by repeated execution:

$$T_b(n) = a_b n + b_b,$$

where b_b is the fixed cost of compilation and initial overhead, and a_b is the average latency, allowing us to infer the inflection point for different backends as a function of the number of repetitions n . Equivalently, writing $b_b = T_C$ (compilation time) and $a_b = T_X$ (per-execution latency), the model reads $T_b(n) = T_C + n T_X$; this is the framework we use throughout to compare backends.

The model $T_b(n) = a_b n + b_b$, together with its *lower envelope*—the curve that, for each value of n , selects the smallest cost among all backends—helps explain the crossover points observed between the different *backends*. Thus, the intersections of these functions indicate from which number of executions n each backend becomes more advantageous.

(5) *Applying fold and comparing regimes*. In addition to the standard models, we evaluate a *fold* as offline preprocessing before

TorchInductor graph capture. At inference, the valid ResNet transformation absorbs BatchNorm into the preceding Conv; the conditional Transformer derivation absorbs the affine part of Layer Norm into a subsequent Linear. The reparameterizations are algebraically exact up to floating-point rounding (Appendix A), but require all affected consumers to remain semantically compensated. A LayerNorm value also used by an identity residual edge cannot be absorbed only into one Linear.

Unlike operator fusion, which combines operations during *codegen*, the *fold* reparameterizes the model before graph capture. The fold tables report compiler cost after preprocessing and record transformation time separately; the deployment envelope adds the measured fold time to the folded intercept. Model loading remains outside both intervals, so neither is a full application-startup TTFB. On fixed model copies with nontrivial BatchNorm statistics, the folded outputs pass `allclose(rtol=1e-3, atol=1e-4)` with maximum absolute errors 6.676×10^{-6} (ResNet-18) and 2.441×10^{-4} (ResNet-50), and no BatchNorm modules remain.

For the folded configuration, $T_f(n) = a_f n + b_f$, where the envelope sets b_f to measured fold preprocessing plus the subsequent compiler-cost proxy. The crossover point for which $\Delta T = T_f(n) - T_b(n) = 0$, where $T_b(n)$ is the cost without fold, is

$$n_{\text{cross}} = \frac{b_b - b_f}{a_f - a_b}, \quad \text{provided } a_f \neq a_b.$$

We report a positive value only when the two point-estimate lines cross for $n > 0$. When $b_f < b_b$ and $a_f \leq a_b$, fold dominates for every non-negative n ; the figures use zero as a sentinel for this no-crossover case.

(6) *JSON construction*. For each (*model*, *input*, *backend*), including the models with fold, we generate a JSON that consolidates the benchmark state. It includes the experiment identification, time metrics, fused *kernels*, and other metadata.

4.4 Execution per Compiler and Model

For each ResNet (**backend**, **model**, **input**) cell, we run the pipeline in five independent processes: backend-native model/input preparation outside the compiler timer, compilation under the boundary defined above, 10 *warmups*, $n = 50$ synchronized timed executions, process-level statistics, and export of metrics, raw samples, retries, and IR artifacts.

4.5 Artifact Analysis Script (TVMScript, HLO, Triton)

We implemented a **Python script** that performs textual *parsing* of the program tree and produces an artifact.

5 Evaluation

All timing tables below come from repeated validation campaigns incorporated into this final version.¹

Throughout, *internal units* are call-like units generated by the backend (TIR calls for TVM, Triton wrapper calls for TorchInductor, and HLO fusions for XLA), whereas *external units* are calls delegated

¹Each ResNet, BERT, and GPT-2 backend/input cell uses five isolated cold processes. The previous BERT fold claim was withdrawn after the audit showed that its matcher performed zero rewrites.

to libraries such as cuDNN or cuBLAS. These static IR categories provide structural context but are not a common hardware-kernel metric: an HLO fusion, a TIR function, and a Python-wrapper call need not map one-to-one to dynamic GPU launches.

PyTorch and XLA load models from native libraries; TVM receives an FX import of the PyTorch model. Frontend representation and parameter binding affect the resulting census. We therefore use the IR figures to describe each stack internally and do not interpret a smaller count as proof of a better compiler.




Figure 1: Static call-like units in the optimized ResNet artifacts at input (1, 3, 224, 224). The categories are backend-specific and are not a ranking or a direct count of dynamic GPU launches.

Figure 1 reports the total static census for the two ResNets, while Figures 2 and 3 separate external-library calls from backend-generated units. The decomposition explains implementation structure—TVM generates most units itself, while TorchInductor and XLA delegate convolutions—but cannot establish fusion quality across unlike IRs. A normalized fusion comparison would require a shared input graph, parameter contract, and counting level.




Figure 2: External-library call-like units in the optimized ResNet artifacts at input (1, 3, 224, 224).

5.1 Residual Networks

In ResNet-18/50, the three native stacks reveal structural differences in generated units and external-library delegation. Table 1 is generated directly from the complete repeated-run JSON grid; the artifact records the synchronization, cache-isolation, retry, and version metadata used for every cell.

Table 1: ResNet-18 and ResNet-50 native-stack compilation-cost proxy and execution time (mean $\pm$ sample sd over $K = 5$ independent process means, ms); each process uses 10 warmups and 50 synchronized executions. The lowest observed mean is marked in bold only when its marginal 95% process-level CI does not overlap either competitor. Compilation boundaries and native model contracts differ as described in Section 4.

Framework	Input (N,C,H,W)	ResNet-18		ResNet-50	
		Comp.	Exec.	Comp.	Exec.
TorchInductor	(1, 3, 224, 224)	11765 $\pm$ 1036	1.88 $\pm$ 0.03	24424 $\pm$ 1539	3.78 $\pm$ 0.08
	(16, 3, 224, 224)	14958 $\pm$ 907	11.86 $\pm$ 0.16	22519 $\pm$ 1094	31.09 $\pm$ 0.11
	(64, 3, 224, 224)	13116 $\pm$ 791	42.43 $\pm$ 0.13	21689 $\pm$ 586	112.51 $\pm$ 0.24
	(1, 3, 512, 512)	15544 $\pm$ 1364	4.85 $\pm$ 0.08	23221 $\pm$ 1411	11.85 $\pm$ 0.09
	(1, 3, 1024, 1024)	16502 $\pm$ 1110	15.03 $\pm$ 0.17	23834 $\pm$ 548	40.19 $\pm$ 0.14
TVM	(1, 3, 224, 224)	8104 $\pm$ 79	8.26 $\pm$ 0.07	13978 $\pm$ 28	20.01 $\pm$ 0.07
	(16, 3, 224, 224)	8333 $\pm$ 42	97.56 $\pm$ 0.06	14318 $\pm$ 130	270.32 $\pm$ 0.54
	(64, 3, 224, 224)	8216 $\pm$ 36	382.88 $\pm$ 1.19	14184 $\pm$ 47	1063.30 $\pm$ 1.27
	(1, 3, 512, 512)	7734 $\pm$ 34	33.01 $\pm$ 0.13	13564 $\pm$ 74	97.22 $\pm$ 0.22
	(1, 3, 1024, 1024)	7727 $\pm$ 44	128.08 $\pm$ 0.31	13577 $\pm$ 107	383.86 $\pm$ 0.37
XLA	(1, 3, 224, 224)	1829 $\pm$ 12	2.13 $\pm$ 0.04	3503 $\pm$ 85	4.19 $\pm$ 0.05
	(16, 3, 224, 224)	2408 $\pm$ 27	11.61 $\pm$ 0.05	4800 $\pm$ 33	29.29 $\pm$ 0.13
	(64, 3, 224, 224)	4560 $\pm$ 100	38.71 $\pm$ 0.12	9352 $\pm$ 121	105.22 $\pm$ 0.26
	(1, 3, 512, 512)	2599 $\pm$ 54	4.23 $\pm$ 0.02	4013 $\pm$ 27	10.38 $\pm$ 0.04
	(1, 3, 1024, 1024)	2663 $\pm$ 76	12.94 $\pm$ 0.01	5011 $\pm$ 73	34.54 $\pm$ 0.02




Figure 3: Backend-generated call-like units in the optimized ResNet artifacts at input (1, 3, 224, 224).

TVM exhibits the largest static call-like census at the representative input. In its imported representation, convolutions and batch normalizations remain separate TIR functions in many blocks, while only shorter element-wise chains are combined. This is a property of the evaluated import and lowering path, not a universal limitation of TVM.

TorchInductor’s inspected wrapper delegates convolutions to external calls and generates surrounding *BatchNorm+ReLU* work with Triton. Moving from ResNet-18 to ResNet-50 increases the generated code and static wrapper-call census, which is associated with a larger compilation-cost proxy (Table 1).

In XLA’s optimized HLO, convolution, bias, normalization, and activation appear in cuDNN custom calls, yielding a smaller HLO-level census. XLA also closes model variables over the JIT function, unlike TorchInductor’s dynamic module-parameter contract. The timing difference is therefore a native-stack result and cannot be attributed solely to fusion or pass ordering.

Table 2: BERT — compilation-cost proxy and execution time (mean $\pm$ sample sd over $K = 5$ independent process means, ms); each process uses 10 warmups and 50 synchronized executions.

Framework	Input	Comp.	Exec.
TorchInductor	(1, 64)	9076 $\pm$ 259	4.95 $\pm$ 0.06
	(1, 128)	9033 $\pm$ 262	8.07 $\pm$ 0.06
	(8, 128)	10466 $\pm$ 759	41.76 $\pm$ 0.03
TVM	(1, 64)	8120 $\pm$ 527	10.83 $\pm$ 0.12
	(1, 128)	7892 $\pm$ 48	22.25 $\pm$ 0.06
	(8, 128)	8259 $\pm$ 69	134.37 $\pm$ 0.35
XLA	(1, 64)	4951 $\pm$ 75	3.56 $\pm$ 0.06
	(1, 128)	5262 $\pm$ 110	5.19 $\pm$ 0.06
	(8, 128)	5730 $\pm$ 34	29.88 $\pm$ 0.07

5.2 Attention-Based Architectures (BERT)

The **BERT** model has a distinct fusion pattern that reflects its architecture centered on *matmuls* [9], QKV (*query-key-value*) projections, and normalization and activation operations.

TVM’s current BERT artifacts contain a larger static `call_tir` census at the imported Relax/TIR level. Its observed compilation means are 7.9–8.3 s, while execution grows from 10.83 ms at (1, 64) to 134.37 ms at (8, 128). These are properties of the evaluated import and lowering path, not universal TVM behavior.

TorchInductor delegates matrix multiplications to external libraries such as cuBLAS [1] and generates surrounding Triton code. Its BERT compilation means are 9.0–10.5 s and its execution means are 4.95–41.76 ms. Generated-code summaries are identical across the five replicas of each cell, but this structural stability does not establish that a particular template caused the timing difference.

XLA’s optimized HLO contains both fusion computations and custom calls. It has the lowest observed compilation and execution means in all three BERT inputs: 5.0–5.7 s and 3.56–29.88 ms, respectively. This ordering is a result for the measured native stacks;

Table 3: GPT-2 — compilation-cost proxy and execution time (mean $\pm$ sample sd over $K = 5$ independent process means, ms); each process uses 10 warmups and 50 synchronized executions.

Framework	Input	Comp.	Exec.
TorchInductor	(1, 16)	9389 $\pm$ 510	3.20 $\pm$ 0.11
	(1, 128)	8939 $\pm$ 187	7.29 $\pm$ 0.04
	(8, 128)	9677 $\pm$ 242	39.79 $\pm$ 0.15
	(8, 256)	9621 $\pm$ 148	82.78 $\pm$ 0.10
TVM	(1, 16)	6720 $\pm$ 250	5.55 $\pm$ 0.13
	(1, 128)	6687 $\pm$ 281	23.69 $\pm$ 0.37
	(8, 128)	6907 $\pm$ 228	149.13 $\pm$ 1.29
	(8, 256)	6540 $\pm$ 14	415.13 $\pm$ 0.34
XLA	(1, 16)	3453 $\pm$ 24	2.76 $\pm$ 0.07
	(1, 128)	4441 $\pm$ 128	4.92 $\pm$ 0.04
	(8, 128)	5860 $\pm$ 102	29.64 $\pm$ 0.21
	(8, 256)	5722 $\pm$ 58	59.67 $\pm$ 0.04

different frontend representations and parameter contracts prevent attributing it solely to fusion.

The process, rather than an individual timed execution, is the experimental unit. Table 2 reports variability over the five independently compiled process means; all 250 per-iteration samples remain available in the artifact.

In short, the BERT model highlights how architectural differences affect each compiler’s strategy.

5.3 Autoregressive Attention-Based Models (GPT-2)

The GPT-2 blocks follow a relatively linear decoder-only pattern—autoregressive attention and MLP, each preceded by LayerNorm and followed by GELU activation, repeated across all layers [10].

TVM’s imported GPT-2 IR has a larger static call_tir census than its BERT import. TVM compilation means remain between 6.5 and 6.9 s, while its execution mean rises from 5.55 ms at (1, 16) to 415.13 ms at (8, 256). The fixed-shape import is executable for all four inputs after the export normalization described in Section 4.

TorchInductor similarly delegates matrix multiplications and generates surrounding Triton code. Its compilation means are 8.9–9.7 s and execution means are 3.20–82.78 ms. The five replicas in each cell produce identical structural summaries, but the experiment does not isolate which generated component explains the timing.

XLA combines HLO fusions with external-library integration. It has the lowest observed compilation and execution means in all four GPT-2 inputs: 3.5–5.9 s and 2.76–59.67 ms, respectively. These process-level results confirm the ordering on this software/hardware stack without establishing a universal compiler ranking.

The current artifacts show different backend-specific structures for TorchInductor, TVM, and XLA, but their counting levels are not normalized and do not establish that fusion structure caused the observed GPT-2 timings.

Across BERT and GPT-2, all 21 backend/input cells contain five independent cold-process observations and 250 synchronized execution samples. They support process-level uncertainty estimates

for the measured stacks, while causal fusion claims still require a shared frontend and controlled ablations.

5.4 Fold

Table 4: ResNet-18 (TorchInductor) — effect of folding BN into Conv. Compilation is mean $\pm$ sample sd across $K = 5$ cold processes; execution entries are pooled means (ms).

Input	Comp. w/o fold	Comp. w/ fold	Δ Comp.	Exec. w/o fold	Exec. w/ fold	Δ Exec.
(1,3,224,224)	11765 $\pm$ 1036	6929 $\pm$ 235	–41.1%	1.88 $\pm$ 0.03	1.86 $\pm$ 0.03	–0.9%
(16,3,224,224)	14958 $\pm$ 907	9026 $\pm$ 214	–39.7%	11.86 $\pm$ 0.16	11.79 $\pm$ 0.10	–0.6%
(64,3,224,224)	13116 $\pm$ 791	9246 $\pm$ 215	–29.5%	42.43 $\pm$ 0.13	42.26 $\pm$ 0.14	–0.4%
(1,3,512,512)	15544 $\pm$ 1364	10201 $\pm$ 1132	–34.4%	4.85 $\pm$ 0.08	4.80 $\pm$ 0.10	–1.0%
(1,3,1024,1024)	16502 $\pm$ 1110	10187 $\pm$ 328	–38.3%	15.03 $\pm$ 0.17	14.89 $\pm$ 0.13	–0.9%

Table 5: ResNet-50 (TorchInductor) — effect of folding BN into Conv. Compilation is mean $\pm$ sample sd across $K = 5$ cold processes; execution entries are pooled means (ms).

Input	Comp. w/o fold	Comp. w/ fold	Δ Comp.	Exec. w/o fold	Exec. w/ fold	Δ Exec.
(1,3,224,224)	24424 $\pm$ 1539	17427 $\pm$ 661	–28.6%	3.78 $\pm$ 0.08	3.71 $\pm$ 0.13	–1.8%
(16,3,224,224)	22519 $\pm$ 1094	19700 $\pm$ 475	–12.5%	31.09 $\pm$ 0.11	30.63 $\pm$ 0.14	–1.5%
(64,3,224,224)	21689 $\pm$ 586	21280 $\pm$ 188	–1.9%	112.51 $\pm$ 0.24	110.15 $\pm$ 0.08	–2.1%
(1,3,512,512)	23221 $\pm$ 1411	19435 $\pm$ 295	–16.3%	11.85 $\pm$ 0.09	11.39 $\pm$ 0.10	–3.9%
(1,3,1024,1024)	23834 $\pm$ 548	21031 $\pm$ 273	–11.8%	40.19 $\pm$ 0.14	38.91 $\pm$ 0.13	–3.2%

Fold (ResNets). Applying the offline fold reduces the mean compiler-cost proxy in all ten ResNet inputs (Tables 4 and 5): 29.5%–41.1% for ResNet-18 and 1.9%–28.6% for ResNet-50. The measured transformation itself takes 100–106 ms for ResNet-18 and 114–120 ms for ResNet-50. Two-sided Welch tests over the five independent process means, with Holm correction across the ten inputs, resolve nine compiler-only differences and the same nine fixed-cost differences after adding preprocessing. The exception is ResNet-50 at (64, 3, 224, 224), whose base compilation sample has high variance. Execution effects are smaller: no ResNet-18 difference is resolved, whereas ResNet-50 has Holm-resolved speedups at (16, 3, 224, 224), (64, 3, 224, 224), (1, 3, 512, 512), and (1, 3, 1024, 1024). These tests compare independently compiled variants; the 250 within-process timing samples are not treated as independent compilations.

The structural evidence explains the compilation result without implying fewer runtime launches. In ResNet-18, the generated wrapper retains 18 Triton and 21 external calls after folding, while generated-code size falls by 38–44% across repetitions. In ResNet-50, wrapper calls remain 50 Triton and 54 external calls, the number of distinct named Triton kernels falls from 21 to 19, and code size falls by about 37%. Fifteen additional anonymous ResNet-50 autotuning stubs are excluded from runtime-call counts. Thus, the supported mechanism is reduced *codegen* volume; attributing the variation across inputs to a particular compiler subsystem would require an additional ablation.

Fold applicability (BERT). Equation (X) in Appendix A is exact for a true single-consumer LayerNorm→Linear chain. The standard Hugging Face BertModel used here is instead post-LayerNorm: each encoder LayerNorm output is consumed both by subsequent

Table 6: BERT (TorchInductor) fold-applicability audit. Base and fold-control are identical because the post-LayerNorm residual graph has no safe LayerNorm→Linear fold. Compilation is mean $\pm$ sample sd across $K = 5$ isolated cold processes; execution is pooled mean $\pm$ sd (ms).

Input	Comp. base	Comp. control	Δ Comp.	Exec. base	Exec. control	Δ Exec.
(1,64)	9422 $\pm$ 217	9373 $\pm$ 409	−0.5%	3.335 $\pm$ 0.102	3.399 $\pm$ 0.157	+1.9%
(1,128)	9405 $\pm$ 173	9351 $\pm$ 228	−0.6%	5.567 $\pm$ 0.316	5.542 $\pm$ 0.328	−0.5%
(8,128)	10311 $\pm$ 518	10396 $\pm$ 590	+0.8%	29.393 $\pm$ 0.859	29.774 $\pm$ 0.531	+1.3%

projections and by an identity residual path. Absorbing (γ, β) only into the projections would therefore change the residual and the model’s inference function. No legal pair is present.

We audited the original BERT scripts and re-ran a conservative fold-control that deep-copies the same graph but performs zero rewrites. The structural matcher finds zero adjacent LayerNorm→Linear siblings in this model; nevertheless, the superseded harness compiled the base and the unchanged “folded” copy sequentially in one process, allowing the second compilation to reuse compiler state. In the corrected experiment, base and control reproduce the submitted BERT model/input protocol but are compiled in distinct cold processes with identical seeds. All five artifacts per variant have the same structural census and code size; the 95% CIs overlap for both compilation and execution in every input (Table 6). The invalid earlier timing claim is therefore withdrawn and no value from that run is used in this paper. A future BERT fold experiment must target an architecture with a genuine single-consumer LayerNorm→Linear edge or explicitly compensate every residual consumer.

In summary, the artifacts expose different backend structures: TVM generates most ResNet call-like units internally, TorchInductor combines Triton wrappers with external-library calls, and XLA combines HLO fusions with custom calls. Because the frontend graphs and counting levels differ, these observations are descriptive rather than a fusion-quality ranking.

TorchInductor delegates convolutions and generates surrounding element-wise work through Triton. In the inspected XLA HLO, convolution, bias, normalization, and activation can appear in a cuDNN custom call. This difference is associated with the measured profiles but is confounded by native frontend and parameter-binding differences.

The valid ResNet tests show that offline *Conv–BatchNorm* preprocessing lowers TorchInductor compilation cost after preprocessing. This demonstrates a guarded frozen-weight specialization opportunity for the evaluated inference path, not that TorchInductor’s internal passes execute in an incorrect order. Figure 4 summarizes the point-estimate deployment regimes.

Execution-time effects are small and uncertain in ResNet-18, showing that **longer compilation does not, by itself, imply lower latency**. In ResNet-50, the folded point estimate is lower in every input, with four differences remaining resolved after correction for multiple comparisons. This motivates reporting both fixed cost and steady-state latency.

Figure 4 reports the point-estimate crossover n_{cross} (Section 4) between folded and original TorchInductor. “Fold dominates” means

that the folded variant has both a lower mean compilation intercept and no higher mean execution slope. A positive entry has the opposite direction from a conventional “optimization pays after” threshold: fold is preferable from $n = 0$ until the displayed value, and the original is preferable afterward. No input yields such an entry: in all ten cells the folded variant satisfies both conditions, so under the point estimates it dominates the original from $n = 0$ onward.

Point-estimate regime: folded vs. original TorchInductor

	ResNet-18	ResNet-50
(1,3,224,224)	fold dominates	fold dominates
(16,3,224,224)	fold dominates	fold dominates
(64,3,224,224)	fold dominates	fold dominates
(1,3,512,512)	fold dominates	fold dominates
(1,3,1024,1024)	fold dominates	fold dominates

Figure 4: Point-estimate regime between folded and original TorchInductor. “Fold dominates” means lower mean fixed cost (including preprocessing) and no higher mean execution; “fold until n ” means that the original point-estimate line becomes lower after n . No input yields a finite boundary under the $K = 5$ point estimates.

Where such a boundary appears, it marks where the preferred measured configuration changes, and its direction depends on both intercept and slope: an optimization with a higher fixed cost but lower latency pays after the crossover, whereas an offline fold with a lower post-preprocessing intercept but slightly higher latency pays before it. For the fold measured here neither case arises, since it is preferable in both terms; such point estimates can guide deployment only when the slope difference is sufficiently stable.

We performed the analysis of the fusion pass in isolation, but the proposed benchmark can be extended to other passes and other models.

5.5 Proposed Benchmark

We developed a *benchmark* in which the user chooses the **model input** ($N \times C \times H \times W$) and the system measures the backend-specific compilation-cost proxy and synchronized inference latency. It reports (i) the measured native stack with the lowest point-estimate total cost for a requested n , and (ii) the corresponding ranges. These recommendations are conditional on the recorded machine, versions, flags, and native model contracts.

The benchmark standardizes the outer sampling conditions—named ResNet architecture, input shape, FP32 precision, GPU, inference mode, $K = 5$ isolated cold processes, 10 warmups, 50 synchronized iterations, and JSON serialization. It deliberately retains each backend’s native model, layout, parameter binding, and compilation API. The resulting recommendation therefore applies to the measured software stacks; it must not be read as a controlled compiler-only comparison.

We model the total cost of each *backend* b as $T_b(n) = a_b n + b_b$, where a_b is the **average time per inference** and b_b is the **compilation time**. By constructing the **lower envelope** of these lines, we obtain the **crossover points** and, consequently, the **dominance ranges** of each compiler.

Report (JSON) and quick reading. The *benchmark* generates a JSON report with the experiment context, raw process/timing measurements, failed XLA attempts, IR metadata, and the point-estimate $T_b(n)$ recommendation. The artifact documentation gives the serialized field names.




Figure 5: Example benchmark plot using the $K = 5$ point estimates: lines $T_b(n) = a_b n + b_b$ per measured configuration, with the lower envelope and crossover points. For ResNet-18 at $(1, 3, 224, 224)$, XLA has the lowest point-estimate cost up to $n \approx 19,378$ and folded TorchInductor thereafter; TVM and the original TorchInductor variant never lie on this lower envelope. The folded intercept includes measured preprocessing.

How to read Figure 5. Each line starts at the backend-specific fixed-cost proxy b_b , and its slope is the measured per-execution latency a_b . The lower envelope highlights the lowest point estimate in each range of n ; dashed lines mark crossovers. It is an illustration of the measured stacks, not an uncertainty-aware compiler ranking.

6 Conclusion

This work applied a repeated cold-process protocol to native TorchInductor, XLA, and TVM stacks on ResNet-18/50, BERT, and GPT-2, and used the lower envelope of $T_b(n) = a_b n + b_b$ as a deployment-oriented summary. The five-process estimates show that the preferred measured stack can change with n : fixed compilation cost matters at small n , whereas the smallest measured steady-state slope determines sufficiently large n . Static fusion or call count alone does not predict that ordering. The cross-backend results characterize complete native stacks and do not isolate compiler quality because frontend graphs, layouts, parameter binding, and timed compilation boundaries differ.

Limitations. Although five independent cold compilations allow process-level uncertainty estimates, $K = 5$ remains small and all measurements come from one machine. An environment control found intermittent XLA convolution-plan failures with the Python

environment’s packaged cuDNN 9.10.2; the reported ResNet campaign instead loads the host cuDNN 9.11.0, for which the repeated grid records any remaining failed attempt. For the 1024^2 inputs, XLA’s autotuner also rejects candidate plans whose workspaces exceed the 8-GB GPU, so its selected plan is hardware-memory constrained. The Transformer campaign uses random initial weights and fixed shapes in the archived container; its absolute values therefore apply to that recorded stack. Future work should increase K , diversify hardware/software stacks, use trained checkpoints, and align frontend graphs and compilation boundaries.

Another limitation lies in the linearity assumption of the model $T_b(n) = a_b n + b_b$ itself: by assuming that the per-batch cost is approximately constant and that all of the fixed cost is paid in the first batch, the model abstracts away cache effects, *autotuning* strategies, and possible recompilations triggered by dynamic *shapes*, which may make the effective behavior closer to a piecewise model than to a single line. In practice, the coefficients a_b and b_b should be read as summary statistics of the range of n that was measured, and not as exact predictions for any usage regime. Future work can refine this dimension by reporting goodness-of-fit metrics (for example, the regression’s R^2) and by investigating nonlinear or *piecewise* models in scenarios with greater temporal variability.

Beyond the comparison between backends, this work also evaluated *fold* transformations in the **TorchInductor** graph, aiming to reduce the compilation cost without changing the model’s semantics. In convolutional architectures (ResNet-18/50), folding BatchNorm into Conv reparameterizes the convolution’s weights and biases and removes the explicit BatchNorm from the graph in *eval* mode. This does not turn the Conv+BatchNorm+ReLU sequence into a single kernel—unlike XLA, which often delegates the whole chain to a single *custom call*—but it reduces the volume of *codegen*, since reparameterizing the Conv is much cheaper than generating a separate Triton kernel for BN+ReLU.

For attention-based models, the analogous LayerNorm+Linear reparameterization is valid only when every use of the affine LayerNorm output can be absorbed or compensated. The post-LayerNorm BERT encoder evaluated here violates this precondition because the same value also follows an identity residual edge. Its isolated control experiment consequently performs no rewrite and shows no resolved timing difference. This negative result narrows the optimization opportunity: an automatic compiler pass must prove the consumer and residual conditions rather than matching only adjacent operator names.

A key observation is that the evaluated TorchInductor path does not perform the valid ResNet fold automatically under its current parameter contract. Offline folding produces a Holm-resolved fixed-cost reduction in nine of ten ResNet inputs and Holm-resolved execution gains in four ResNet-50 inputs. This motivates guarded frozen-weight specialization with explicit legality checks, while leaving open whether automatic specialization is profitable under other contracts.

As a direct next step, a compiler pass could (i) specialize frozen BatchNorm parameters into Conv in `eval()` mode, (ii) request a convolution–activation epilogue when supported, and (iii) fall back safely otherwise. Its evaluation must include transformation overhead and dynamic launch profiling; the present result establishes lower code-generation volume, not fewer runtime launches.

Artifact Availability

The artifact provides the benchmark source code, pinned Python environments, a GPU-enabled Docker image recipe, correctness and smoke tests, complete repeated cold-process ResNet and Transformer JSON grids, documented compiler-IR outputs, and scripts that regenerate the statistical analyses, tables, and plots. It is archived under the MIT License at <https://doi.org/10.5281/zenodo.21731237>, which will resolve to the specific deposited version; the development repository is mirrored at <https://github.com/kameczera/compiler-benchmark>. The README documents the hardware and software requirements, expected outputs, validation procedure, resource costs, and the commands needed for both a short functional check and the complete experiment.

Acknowledgements

The authors used **ChatGPT (OpenAI, 2025)** as a support tool during the preparation of this paper, restricted to the following purposes: refining the writing and checking textual coherence; reviewing the algebraic derivations of the kernel fusions in Appendix A; and assisting in writing and debugging portions of the benchmark code used in the experiments. All technical, methodological, and analytical content is the authors' own and was reviewed and validated by the authors. This disclosure follows the SBLP generative-AI policy and the corresponding ACM, IEEE, and Springer guidelines.

A Derivation of the Analyzed Fusion and Fold Transformations

This appendix distinguishes one kernel-fusion expression from two conditional parameter-folding derivations. Throughout, $\odot$ denotes the element-wise (Hadamard) product, broadcast along the channel dimension when the operands' shapes differ.

BatchNorm + ReLU

This widely known fusion combines a batch normalization (BatchNorm) with a following rectified linear unit (ReLU). Let $X \in \mathbb{R}^{N \times C \times H \times W}$ be the activation tensor (batch, channels, spatial dimensions); $\mu, \sigma^2 \in \mathbb{R}^C$ the per-channel estimated mean and variance; $\gamma, \beta \in \mathbb{R}^C$ the learnable scale and shift; and $\varepsilon > 0$ the numerical-stability constant. BatchNorm computes

$$(I) \quad Y = \frac{X - \mu}{\sqrt{\sigma^2 + \varepsilon}} \odot \gamma + \beta. \quad (1)$$

The ReLU is then applied element-wise to the output:

$$(II) \quad Z = \max(0, Y). \quad (2)$$

An unfused realization may launch (I) and (II) separately, whereas a fused implementation computes both in one kernel:

$$(III) \quad y = \max\left(0, \gamma \cdot \frac{x - \mu}{\sqrt{\sigma^2 + \varepsilon}} + \beta\right). \quad (3)$$

Modern deep-learning compilers cannot fuse every such pair profitably or legally, for the reasons discussed in Section 3.

Conv + BatchNorm + ReLU

In computer vision, the Conv--BatchNorm--ReLU chain contains two distinct opportunities: at inference, a fixed BatchNorm can be folded into the preceding convolution parameters ($\tilde{W}, \tilde{b}$), and, separately, a backend may implement ReLU as a convolution epilogue. The evaluated transformation performs the parameter fold only.

Formally, let $X \in \mathbb{R}^{N \times C_{in} \times H \times W}$ be the input and $W \in \mathbb{R}^{C_{out} \times C_{in} \times K_H \times K_W}$, $b \in \mathbb{R}^{C_{out}}$ the convolution parameters. The pre-activation output is

$$(IV) \quad U_{n,c,h,w} = \sum_{c'=1}^{C_{in}} \sum_{r,s} W_{c,c',r,s} X_{n,c',h+r,w+s} + b_c. \quad (4)$$

BatchNorm at inference, with parameters $\mu, \sigma^2, \gamma, \beta \in \mathbb{R}^{C_{out}}$ and $\varepsilon > 0$, is

$$(V) \quad Y_{n,c,h,w} = \frac{U_{n,c,h,w} - \mu_c}{\sqrt{\sigma_c^2 + \varepsilon}} \gamma_c + \beta_c. \quad (5)$$

The ReLU (II) is then applied element-wise. We absorb the BatchNorm into the convolution parameters by defining

$$(VI) \quad \tilde{W}_{c,c',r,s} = \frac{\gamma_c}{\sqrt{\sigma_c^2 + \varepsilon}} W_{c,c',r,s}, \quad (6)$$

$$(VII) \quad \tilde{b}_c = \frac{\gamma_c}{\sqrt{\sigma_c^2 + \varepsilon}} (b_c - \mu_c) + \beta_c. \quad (7)$$

Substituting into Y and applying ReLU yields an equivalent expression that a backend may implement as one convolution-activation operation:

$$(VIII) \quad Z_{n,c,h,w} = \max\left(0, \sum_{c'=1}^{C_{in}} \sum_{r,s} \tilde{W}_{c,c',r,s} X_{n,c',h+r,w+s} + \tilde{b}_c\right). \quad (8)$$

The benchmark guarantees only the reparameterization by $(\tilde{W}, \tilde{b})$; whether Equation (VIII) becomes one runtime kernel is backend-dependent.

LayerNorm (Affine Part) + Linear

Conditionally, the affine part of LayerNorm (LN) can be reparameterized into every subsequent Linear consumer. For an activation vector $h \in \mathbb{R}^d$, write the LN as $\tilde{h} = \gamma \odot \hat{h} + \beta$, where $\hat{h} = \text{Norm}(h)$ is centered and normalized by the mean and variance of h , and $\gamma, \beta \in \mathbb{R}^d$ are the learnable scale and shift. For a Linear layer with weights $W \in \mathbb{R}^{m \times d}$ and bias $b \in \mathbb{R}^m$, expanding $y = W\tilde{h} + b = W(\gamma \odot \hat{h} + \beta) + b$ gives

$$(IX) \quad y = W \text{diag}(\gamma) \hat{h} + (W\beta + b). \quad (9)$$

Defining $W' = W \text{diag}(\gamma)$ and $b' = W\beta + b$, the LN+Linear composition becomes

$$(X) \quad y = W' \text{Norm}(h) + b'. \quad (10)$$

Absorbing (γ, β) is thus legal only if every use of the affine output is removed or compensated, which the post-LayerNorm BERT evaluated here violates through its residual uses: Equation (X) is a conditional derivation, not a BERT optimization result.

References

- [1] Jason Ansel, Edward Yang, Horace He, Natalia Gimelshein, Animesh Jain, Michael Voznesensky, Bin Bao, Peter Bell, David Berard, Evgeni Burovski, et al. 2024. PyTorch 2: Faster Machine Learning Through Dynamic Python Bytecode Transformation and Graph Compilation. In *Proceedings of the 29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 2*. ACM, 929–947. doi:10.1145/3620665.3640366
- [2] Xuyi Cai, Ying Wang, and Lei Zhang. 2022. Optimus: An Operator Fusion Framework for Deep Neural Networks. *ACM Transactions on Embedded Computing Systems* 22, 1 (2022), 26 pages. doi:10.1145/3520142
- [3] Michael Canesche, Vanderson Martins do Rosario, Edson Borin, and Fernando Magno Quintão Pereira. 2025. Fusion of Operators of Computational Graphs via Greedy Clustering: The XNNC Experience. In *Proceedings of the 34th ACM SIGPLAN International Conference on Compiler Construction*. ACM, 117–127. doi:10.1145/3708493.3712689
- [4] Tianqi Chen, Thierry Moreau, Ziheng Jiang, Lianmin Zheng, Eddie Yan, Haichen Shen, Meghan Cowan, Leyuan Wang, Yuwei Hu, Luis Ceze, Carlos Guestrin, and Arvind Krishnamurthy. 2018. TVM: An Automated End-to-End Optimizing Compiler for Deep Learning. In *13th USENIX Symposium on Operating Systems Design and Implementation (OSDI 18)*. 578–594.
- [5] Alain Darté. 1999. On the complexity of loop fusion. In *Proceedings of the 1999 International Conference on Parallel Architectures and Compilation Techniques*. IEEE, 149–157. Cat. No. PR00425.
- [6] Christopher W. Fraser. 1979. A Compact, Machine-Independent Peephole Optimizer. In *Proceedings of the 6th ACM SIGACT-SIGPLAN Symposium on Principles of Programming Languages (POPL)*. ACM, 1–6. doi:10.1145/567752.567753
- [7] Chris Leary and Todd Wang. 2017. XLA: TensorFlow, Compiled. In *TensorFlow Dev Summit*. <https://www.tensorflow.org/xla>
- [8] Wei Niu, Jiexiong Guan, Yanzhi Wang, Gagan Agrawal, and Bin Ren. 2021. DNNFusion: Accelerating Deep Neural Networks Execution with Advanced Operator Fusion. In *Proceedings of the 42nd ACM SIGPLAN International Conference on Programming Language Design and Implementation*. ACM, 883–898. doi:10.1145/3453483.3454083
- [9] Suchita Pati, Shaizeen Aga, Nuwan Jayasena, and Matthew D. Sinclair. 2021. Demystifying BERT: Implications for Accelerator Design. In *IEEE International Symposium on Workload Characterization (IISWC 2021)*. IEEE, 109–120. doi:10.1109/IISWC53511.2021.00019
- [10] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. 2019. Language Models are Unsupervised Multitask Learners. *OpenAI Technical Report* (2019).
- [11] Daniel Snider and Ruofan Liang. 2023. Operator Fusion in XLA: Analysis and Evaluation. <https://arxiv.org/abs/2301.13062>. arXiv:2301.13062.
- [12] Michael E. Wolf and Monica S. Lam. 1991. A Data Locality Optimizing Algorithm. In *ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*. 30–44.